

## CS 343 Winter 2018 – Assignment 2

Instructor: Peter Buhr


**Due Date: Wednesday, January 31, 2018 at 22:00**

**Late Date: Friday, February 2, 2018 at 22:00**

January 16, 2018

This assignment has complex semi-coroutines, and introduces full-coroutines and concurrency in  $\mu\text{C++}$ . Use it to become familiar with these new facilities, and ensure you use these concepts in your assignment solution, i.e., writing a C-style solution for questions is unacceptable, and will receive little or no marks. (You may freely use the code from these [example programs](#).)

1. Write a program that filters a stream of text. The filter semantics are specified by command-line options. The program creates a *semi-coroutine* filter for each command-line option joined together in a pipeline with a reader filter at the input end of the pipeline, followed by the command-line filters in the middle of the pipeline (maybe zero), and a writer filter at the output end of the pipeline. Control passes from the reader, through the filters, and to the writer, ultimately returning back to the reader. One character moves along the pipeline at any time. For example, a character starts from the reader filter, may be deleted or transformed by the command-line filters, and any character reaching the writer filter is printed.



(In the following, you may not add, change or remove prototypes or given members; you may add a destructor and/or private and protected members.)

Each filter must inherit from the abstract class Filter:

```
_Coroutine Filter {  
    protected:  
        _Event Eof {};  
        Filter * next;  
        unsigned char ch;  
    public:  
        void put( unsigned char c ) {  
            ch = c;  
            resume();  
        }  
};
```

which ensures each filter has a put routine that can be called to transfer a character along the pipeline.

The reader reads characters from the specified input file and passes these characters to the first coroutine in the filter:

```
_Coroutine Reader : public Filter {  
    // YOU MAY ADD PRIVATE MEMBERS  
    void main();  
    public:  
        Reader( Filter * f, istream * i );  
};
```

The reader constructor is passed the next filter object, which the reader passes one character at a time from the input stream, and an input stream object from which the reader reads characters. No coroutine calls the put routine of the reader; all other coroutines have their put routine called. When the reader reaches end-of-file, it raises the exception Eof at the next filter, resumes the next filter with an arbitrary character, and terminates.

The writer is passed characters from the last coroutine in the filter pipeline and writes these characters to the specified output file:

```
_Coroutine Writer : public Filter {
    // YOU MAY ADD PRIVATE MEMBERS
    void main();
public:
    Writer( ostream * o );
};
```

The writer constructor is passed an output stream object to which this filter writes characters that have been filtered along the pipeline. No filter is passed to the writer because it is at the end of the pipeline. When the write receives the Eof exception, it prints out how many characters were printed, e.g.:

16 characters

All other filters have the following interface:

```
_Coroutine filter-name : public Filter {
    // YOU MAY ADD PRIVATE MEMBERS
    void main();
public:
    filter-name( Filter * f, ... );
};
```

Each filter constructor is passed the next filter object, which this filter passes one character at a time after performing its filtering action, and “...” is any additional information needed to perform the filtering action. When a filter reaches end-of-file, it raises the exception Eof at the next filter, resumes the next filter with an arbitrary character, and terminates.

The pipeline is built by the program main from writer to reader, in reverse order to the data flow. Each newly created coroutine is passed to the constructor of its predecessor coroutine in the pipeline. The reader's constructor resumes itself to begin the flow of data, and it calls the put routine of the next filter to begin moving characters through the pipeline to the writer. Normal characters, as well as control characters (e.g., '\n', '\t'), are passed through the pipeline. When the reader reaches end-of-file, it raises the exception Eof at the next filter, resumes the next filter with an arbitrary character, and terminates. Similarly, each coroutine along the filter pipeline raises the exception Eof at the next filter along the pipeline and then terminates. The program main ends when the reader declaration completes, implying all the input characters have been read and all the filter coroutines are terminated. The reader coroutine can read characters one at a time or in groups; the writer coroutine can write characters one at a time or in groups.

Filter options are passed to the program via command line arguments. For each filter option, create the appropriate coroutine and connect it into the pipeline. If no filter options are specified, then the output should simply be an echo of the input from the reader to the writer. *Assume all filter options are correctly specified, i.e., no error checking is required on the filter options.*

The filter options that must be supported are:

- h The *hex dump* option replaces each character in the stream with its corresponding ASCII 2-hexadecimal digit value. For example, the character 'a' is transformed into the two characters '6' and '1', as 61 is the ASCII hexadecimal value for character 'a'. The hexadecimal filter also formats its output by transforming two characters into 4 hexadecimal digits, separated by one space, transforming two characters into 4 hexadecimal digits, separated by three spaces, repeating this sequence four times, and adding a newline. Note, the spaces are separators not terminators, and hence, there are no spaces at the end of a line. For example, this input sequence:

The quick brown fox jumps over the lazy dog.

generates the following:

```

5468 6520    7175 6963    6b20 6272    6f77 6e20
666f 7820    6a75 6d70    7320 6f76    6572 2074
6865 206c    617a 7920    646f 672e    0a

```

Note, it is possible to convert a character to its hexadecimal value using a simple, short expression. No complex library routines are required.

- w** The *whitespace* option removes all spaces and tabs (isblank) from the start and end of lines, and collapses multiple spaces and tabs within a line into a single space. Lines are delimited by the newline character (' \n').

–**you-create** **In addition, design and implement one other *non-trivial* filter operation that might be useful.**

Fully document the new filter in the code; i.e., what it does, how to use it, and show examples of its usage in your testing.

The order in which filter options appear on the command line is significant, i.e., the left-to-right order of the filter options on the command line is the first-to-last order of the filter coroutines. As well, a filter option may appear more than once in a command. Each filter should be constructed without regard to what any other filters do, i.e., there is no communication among filters using global variables; all information is passed using member put. **Hint:** scan the command line left-to-right to locate and remember the position of each option, and then scan the option-position information right-to-left (reverse order) to create the filters with their specified arguments.

The executable program is to be named `filter` and has the following shell interface:

```
filter [ -filter-options ... ] [ infile [outfile] ]
```

(Square brackets indicate optional command line parameters, and do not appear on the actual command line.) If filter options appear, assume they appear *before* the file names. Terminate the program for unknown or insufficient command arguments, or if an input or output file cannot be opened. Assume any given argument values are correctly formed, i.e., no error checking is required for numeric values. If no input file name is specified, input from standard input and output to standard output. If no output file name is specified, output to standard output.

**WARNING:** When writing coroutines, try to reduce or eliminate execution “state” variables and control-flow statements using them. A state variable contains information that is not part of the computation and exclusively used for control-flow purposes (like flag variables). Use of execution state variables in a coroutine usually indicates you are not using the ability of the coroutine to remember prior execution information. *Little or no marks will be given for solutions explicitly managing “state” variables.* See Section 3.1.3 in the [Course Notes](#) for details on this issue. Also, make sure a coroutine’s public methods are used for passing information to the coroutine, but not for doing the coroutine’s work, which must be done in the coroutine’s main.

2. Write a *full coroutine* that simulates the game of Hot Potato. What makes the potato *hot* is that it explodes after its internal count-down timer reaches zero. The game consists of  $N$  players linked in a circle, where one of the players also acts as the umpire. The umpire starts a *game* by randomly tossing the *hot* potato to a player on their left or right. The potato is then randomly tossed left or right among the players until the timer goes off and it explodes. The player holding the potato when the timer goes off (potato explodes) is removed from the game (circle) and deleted. A player cannot delete themselves, so they call the umpire to be deleted. The umpire then resets the timer in the potato, and begins the game again with the remaining players, unless only one player remains, who is the winner.

Should the potato explode for the umpire, a new umpire is *elected* among the remaining players. An election involves traversing the circle to find the remaining player with the highest id; this player becomes the new umpire.

The potato is tossed among the players and it contains the count-down timer that goes off after a random number of clock ticks. The interface for the Potato is (you may only add a public destructor and private members):

```

class Potato {
    // YOU ADD MEMBERS HERE
public:
    _Event Explode {};
    Potato( unsigned int maxTicks = 10 );
    void reset( unsigned int maxTicks = 10 );
    void countdown();
};

```

The constructor is optionally passed the *maximum* number of ticks until the timer goes off. The potato chooses a random value between 1 and the maximum for the number of ticks. Member reset is called by the umpire to re-initialize the timer for the next set. Member countdown is called by the players, and throws exception Explode, if the timer has reached zero. Rather than absolute time to implement the potato's timer, make each call to countdown be one tick of the clock (relative time).

The interface for a Player is (you may only add a public destructor and private members):

```

_Coroutine Player {
    _Event Terminate {
public:
        Player * victim;                // delete player
        Terminate( Player * victim ) : victim( victim ) {}
    };
    _Event Election {
public:
        int id;                        // highest id seen so far
        Election( int id ) : id( id ) {}
    };
    Player * partner[2];                // left and right player
    // YOU ADD MEMBERS HERE
    void main();
public:
    static Player * umpire;              // current umpire

    Player( int id, Potato & potato );
    void start( Player & lp, Player & rp ); // supply partners
    void toss();                          // tossed the potato
    void vote( int id );                  // resume partner to vote
    void terminate();                     // resume umpire to delete
};

```

The exception Terminate is raised at the umpire and contains the player to be deleted. The exception Election is raised at the player on the right to run an election and contains the highest player id seen so far in the circle.

The global variable umpire is the player currently acting as the umpire. This variable allows a player to communicate directly with the umpire. The variable umpire is updated as part of eliminating the umpire from the game and an election selects a new umpire.

The constructor is passed an assigned identification number (0 to  $N - 1$ ) and the potato created by the main program. The start member receives a player's left and right partners. The toss member is called to conceptually pass the potato when a player is *not* eliminated while holding the *hot* potato. A player randomly tosses to the player on their left or right. The vote member is called to vote in an election and is passed the current highest player id seen during traversal of the circle. The terminate member is called to restart the umpire to propagate the terminate exception.

When a player determines it is eliminated, i.e., the timer went off while holding the potato, it starts an election, if it is the umpire. It then unlinks itself from the circle, and raises the Terminate exception at the umpire and resume it via member terminate.

The umpire handles the Terminate exception, deleting the eliminated player. Hence, player coroutines are created in the main program but deleted during the game by the umpire. The umpire then resets the potato, and continues by starting a new game by randomly tossing to its left/right player. This toss counts with respect to the timer in the potato.

The executable program is named `hotpotato` and has the following shell interface:

```
hotpotato [ games | "x" [ players | "x" [ seed | "x" ] ] ]
```

**games** is the number of card games to be played ( $\geq 0$ ). If `x` or no value for `games` is specified, assume 5.

**players** is the number of players in the game ( $\geq 2$ ). If `x` or no value for `players` is specified, generate a random integer in the range from 2 to 10 inclusive for each game.

**seed** is the starting seed for the random number generator to allow reproducible results ( $> 0$ ). If `x` or no value for `seed` is specified, initialize the random number generator with an arbitrary seed value (e.g., `getpid()` or `time`).

Check all command arguments for correct form (integers or `x`) and range; print an appropriate usage message and terminate the program if a value is missing or invalid. The program main creates the potato and all necessary data structures, sets the global umpire to the first player, calls its toss member, and deletes the umpire.

To form the ring of players, the program main calls the start member for each player to link the players together into a circle. The start routine also resumes the player to set the program main as its starter (needed during termination). The main routine of each player suspends back immediately so the next player can be started. After the program main initializes the ring, it sets the global umpire variable to the player with id 0, and tosses the potato to it to start the game.

Generate the following output showing a dynamic display of the game, where U means umpire and E means election:

```
$ hotpotato 1 6 1003
  POTATO goes off after 8 tosses
6 players in the game
U 0 -> 5 -> 0 -> 5 -> 0 -> 5 -> 0 -> 5 is eliminated
  POTATO goes off after 4 tosses
U 0 -> 4 -> 3 -> 2 is eliminated
  POTATO goes off after 1 toss
U 0 is eliminated
E 0 -> 1 -> 3 -> 4 : umpire 4
  POTATO goes off after 7 tosses
U 4 -> 1 -> 4 -> 3 -> 1 -> 4 -> 1 is eliminated
  POTATO goes off after 7 tosses
U 4 -> 3 -> 4 -> 3 -> 4 -> 3 -> 4 is eliminated
E 4 -> 3 : umpire 3
  POTATO goes off after 2 tosses
U 3 wins the Match!
```

To obtain repeatable results, all random numbers are generated using class `PRNG`. There are exactly three calls to the random-number generator: one in `potato::reset`, one in program main, and one in `Player::main`.

**WARNING:** When writing coroutines, try to reduce or eliminate execution “state” variables and control-flow statements using them. A state variable contains information that is not part of the computation and exclusively used for control-flow purposes (like flag variables). Use of execution state variables in a coroutine usually indicates you are not using the ability of the coroutine to remember prior execution information. *Little or no marks will be given for solutions explicitly managing “state” variables.* See Section 3.1.3 in the [Course Notes](#) for details on this issue. Also, make sure a coroutine’s public methods are used for passing information to the coroutine, but not for doing the coroutine’s work, which must be done in the coroutine’s main.

3. Compile the program in Figure 1 using the `u++` command, without and with compilation flag `-multi` and ***no optimization***, to generate a uniprocessor and multiprocessor executable. Run both versions of the program 10 times with command line argument 10000000 on a multi-core computer with at least 2 CPUs (cores).
  - (a) Show the 10 results from each version of the program.
  - (b) Must all 10 runs for each version produce the same result? Explain your answer.
  - (c) In theory, what are the smallest and largest values that could be printed out by this program with an argument of 10000000? Explain your answers. (**Hint:** one of the obvious answers is wrong.)
  - (d) Explain the difference in the size of the values between the uniprocessor and multiprocessor output.

```

#include <iostream>
using namespace std;

volatile int iterations = 10000000, shared = 0; // volatile to prevent dead-code removal

_Task increment {
void main() {
    for ( int i = 1; i <= iterations; i += 1 ) {
        shared += 1;           // two increments to increase pipeline size
        shared += 1;
    }
};
int main( int argc, char * argv[] ) {
    if ( argc == 2 ) iterations = atoi( argv[1] );
#ifdef __U_MULTITHREAD__
    uProcessor p;                // create 2nd kernel thread
#endif // __U_MULTITHREAD__
    {
        increment t[2];
    } // wait for tasks to finish
    cout << "shared:" << shared << endl;
}

```

Figure 1: Interference

## Submission Guidelines

Please follow these guidelines very carefully. Review the [Assignment Guidelines](#) and [C++ Coding Guidelines](#) before starting each assignment. **Each text file, i.e., \*.txt file, must be ASCII text and not exceed 500 lines in length, where a line is a maximum of 120 characters.** Programs should be divided into separate compilation units, i.e., \*.{h,cc,C,cpp} files, where applicable. Use the [submit](#) command to electronically copy the following files to the course account.

1. q1\*.{h,cc,C,cpp} – code for question 1, p. 1. **Program documentation must be present in your submitted code. No user or system documentation is to be submitted for this question. Output for this question is checked via a marking program, so it must match exactly with the given program.**
2. q1\*.testdoc – Documentation for your filter and how to use it, and test documentation for question 1, p. 1, which includes the input and output of your tests. **Poor documentation of how and/or what is tested can result in a loss of all marks allocated to testing.**
3. PRNG.h – random number generator (provided)
4. q2\*.{h,cc,C,cpp} – code for question 1, p. 1. **Program documentation must be present in your submitted code. No user, system or test documentation is to be submitted for this question. Output for this question is checked via a marking program, so it must match exactly with the given program.**
5. q3\*.txt – contains the information required by question 3.
6. Modify the following Makefile to compile the programs for question 1, p. 1 and 2, p. 3 by inserting the object-file names matching your source-file names.

```

CXX = u++                                # compiler
CXXFLAGS = -g -Wall -Wextra -MMD
MAKEFILE_NAME = ${firstword ${MAKEFILE_LIST}} # makefile name

OBJECTS1 = # object files forming 1st executable with prefix "q1"
EXEC1 = filter

OBJECTS2 = # object files forming 2nd executable with prefix "q2"
EXEC2 = hotpotato

```

```

OBJECTS = ${OBJECTS1} ${OBJECTS2}      # all object files
DEPENDS = ${OBJECTS:.o=.d}             # substitute ".o" with ".d"
EXECS = ${EXEC1} ${EXEC2}              # all executables

.PHONY : all clean

all : ${EXECS}                          # build all executables

#####

${EXEC1} : ${OBJECTS1}                  # link step 1st executable
    ${CXX} ${CXXFLAGS} $^ -o $@

${EXEC2} : ${OBJECTS2}                  # link step 2nd executable
    ${CXX} ${CXXFLAGS} $^ -o $@

#####

${OBJECTS} : ${MAKEFILE_NAME}           # OPTIONAL : changes to this file => recompile

-include ${DEPENDS}                     # include *.d files containing program dependences

clean :                                  # remove files that can be regenerated
    rm -f *.d *.o ${EXECS}

```

This makefile is used as follows:

```

$ make filter
$ filter ...
$ make hotpotato
$ hotpotato ...

```

Put this Makefile in the directory with the programs, name the source files as specified above, and then type `make filter` or `make hotpotato` in the directory to compile the programs. This Makefile must be submitted with the assignment to build the program, so it must be correct. Use the web tool [Request Test Compilation](#) to ensure you have submitted the appropriate files, your makefile is correct, and your code compiles in the testing environment.

**Follow these guidelines. Your grade depends on it!**